

Field Intelligence

From a field observation to a verified operational outcome.

I started with one question: **what would make a franchisee trust an AI-assisted finding?** The answer was not a smarter form. It was a field companion that lets consultants work naturally, preserves what they actually observed, and carries every supported issue through human review, ownership, correction and independent verification.

The product in one minute

A consultant walks the property area by area and speaks, types, photographs or records what they notice. Field Intelligence structures that evidence in the background, asks a focused follow-up only when something is unclear, and checks the relevant source-labelled guide. It then prepares a review packet that keeps the consultant's statement separate from the AI interpretation.

A reviewer can approve, edit, challenge, request evidence or reject the packet. An approved issue becomes an assigned case with a due date. The location operator adds before-and-after evidence, and a different person verifies the result. The product turns a moment in the field into a traceable operational outcome.

Three moments that make it feel different

1. Walk naturally

The mobile-first companion supports voice, photo, video, typed notes and guided checks without forcing the consultant to stop and complete a long form. A zone-aware guide keeps the walkthrough focused, while a background capture queue makes progress visible.

2. Make every conclusion defensible

The AI preserves the original statement, shows confidence and uncertainty, cites the guide it used, and states what the evidence does not prove. It asks bounded clarifying questions instead of guessing. Media must match the selected area and reported condition; unrelated photo or video evidence is rejected.

3. Close the loop

Supported issues create a case receipt with an owner, due date and evidence state. Reviewer, operator and verifier responsibilities stay separate. Before-and-after proof, resolution notes, disputes and decisions remain connected in one audit trail.

13 FIELD AREAS	29 GUIDED CHECKS	362 WOLF CREEK REVIEWS
1,235 COMPARATOR REVIEWS	70 REGRESSION TESTS	0/8 UNSUPPORTED GATE

What the working product proves

This is more than a report generator. It is a governed evidence-to-outcome workflow running with live Gemini text, image, audio and video understanding, persistent Postgres data and private evidence storage. Every model call records provider, model, tokens, latency, estimated cost, retries and outcome. A clearly labelled deterministic fixture provider makes every release repeatable.

Confidence is engineered into the workflow

Evidence first

Voice enters analysis only after transcript confirmation. Photos can be recommended or required according to the type of issue. High-privacy areas require an explicit privacy attestation. Evidence is target-linked, content-validated, metadata-scrubbed and served through authenticated routes.

Deliberate AI

The model works through typed, read-only, tenant-scoped tools and schema-validated outputs. Clarification is bounded so conversations converge. A three-lens challenge panel separately tests evidence sufficiency, franchisee fairness and standards fit; disagreement can downgrade or halt a candidate finding.

From insight to operational value


Human governance

AI proposes; people decide. The reviewer controls the finding, the operator owns the correction, and an independent verifier closes it. Exact statements, standard versions, edits, challenges, disputes, evidence and status changes remain traceable.

Customer intelligence with clear evidence boundaries

The product analyzes all **362 anonymized Wolf Creek reviews**, including **42 one-star reviews**, and locally filters recent negative written feedback. It surfaces recurring themes and compares aggregate strengths across **1,235 reviews from three nearby Atlanta courses**. Customer sentiment adds context and opportunity signals; it never becomes proof of a field condition.

One governed operating loop



Built, tested and ready to demonstrate

- **70 automated regression tests** cover authentication, concurrency, model budgets, tenant boundaries, evidence matching, privacy, review governance and verified closure.
- The executable behavioural evaluation passes **14/14 cases**, with live-vision-only cases labelled separately and the **0/8 unsupported-finding gate** clear.
- The production bundle, Docker image and deployed Render service have been exercised across desktop and mobile flows.
- A ready-made showcase audit demonstrates capture, clarification, evidence, review, ticketing, resolution and verification without asking the evaluator to create data first.
- The same engine also runs an **AI Quoz EV & Delivery Depot** tenant with its own locations, zones and EV-specific standards, demonstrating that the workflow is configurable across operating businesses rather than being tied to golf.

A practical 3-6 week field pilot

Week 1 - configure and observe

Load one controlled operating pack, connect company identity, agree the evidence/privacy policy, and shadow a Wolf Creek consultant. Establish baselines for walkthrough time, clarification rate, reviewer edits and dispute outcomes.

Weeks 2-3 - run in shadow mode

Run one property alongside the existing process. Use accepted corrections to refine prompts and classification rules, compare AI packets with human judgment, and measure whether findings are faster to review and easier to defend.

Weeks 4-6 - prove operational value

Connect approved work-order and customer-response channels, operate one live property, and measure closure time, repeat issues and evidence quality. Use the results to decide the next location or brand. The north-star outcome is **faster, fairer resolution that operators and franchisees can trust**.

A five-minute showcase

- 1 **Field Consultant:** open the ready-made visit and inspect the multimodal security case.
- 2 **Reviewer:** see the evidence-to-guide packet, uncertainty and challenge outcome.
- 3 **Location Operator:** compare before and after evidence and record the correction.
- 4 **Brand Leader:** verify the result from the independent closure queue.
- 5 **Technical Evaluator:** open Eval Lab and the model ledger to inspect reliability, cost and traceability.

Field Intelligence shows the product judgment I would bring to BroadPeak: start with the person doing the work, make AI useful without making it authoritative, and design the full operational loop - not only the model call.